



CUDA warp shared-memory: effetto del padding

M=16384, N=16384, P=1, precisione=double, BLOCK=256, granularità=32 · 20 ripetizioni



SMEM_PAD=1 selezionato · prestazione relativa media=99.83% · caso peggiore=99.13% · il padding evita il forte calo osservato a k=32